

Data Structures 1

Wet 2 – dry part

Names of students: Matam Miller, Shahar Fox

Overview of the Data Structures we used:

We use four principal data structures (plus a small “banned ID” hash):

1) **Union-Find** to group teams into “components”:

Each Team node has:

- parent: pointer to another Team or to itself if it is a “root.”
- active: boolean indicating if it is the currently alive root (true) or a loser in a merge (false).
- record: integer sum of all jockey wins minus losses in that root’s entire set.
- teamId: the official ID for that team.

2) **Hash Table** for Teams by teamId.

- Key = integer teamId, Value = pointer to the Team node.
- Used in add_team, findTeamById, merge_teams checks and more.

3) **Hash Table** for Jockeys by jockeyId.

Key = integer jockeyId, Value = pointer to a Jockey struct containing:

- jockeyId
- pointer to the Team node it belongs to
- record: personal jockey win-loss balance

4) **Hash Table** for Teams by record.

- Key = integer (the “record”)
- Value = chain of pointers to Team nodes with that record.
- We primarily store root teams in their record’s chain, enabling quick scans for unite_by_record(record).

5) **Hash Table** for Banned IDs.

- Once a team ID is used and then “loses” in a merge, that ID is “banned” from being added again.
- Checking or inserting in this structure is $O(1)$ expected.

Each hash table uses separate chaining with singly linked lists. We pick a prime size (e.g., 40 001) and a simple mod-based hash function.

Union-Find implementation sketch:

When we merge team A and team B:

- We pick a winner root and a loser root.
- We set loser->active = false and loser->parent = winner.
- The winner’s record is set to be the sum of both records.
- The loser’s ID is inserted into the “banned” structure so that future add_team(loserId) fails.

When we want to find which team a jockey belongs to, we do a find up the parent pointers until we reach an active root. The standard union-find approach ensures near-constant time for repeated merges/finds.

Overview of the functions and their algorithms:

Add_team(teamId):

Algorithm:

1. If teamId ≤ 0 , return INVALID_INPUT.
2. Check the banned hash. If teamId is present, return FAILURE.
3. Look up teamId in the Team hash. If found, return FAILURE.
4. Allocate a new Team object with:
 - parent = self
 - active = true
 - record = 0
 - teamId = teamId
5. Insert it into the Team hash and into the “record = 0” chain in the record hash.
6. Return SUCCESS.

Time Complexity:

- Step 2 (banned hash check) is an expected $O(1)$ hash lookup.
- Step 3 (team hash check) is $O(1)$ expected.

- Steps 5 (insert) is $O(1)$ expected.

Hence the entire function is $O(1)$ expected amortized.

Mathematical Justification:

- A hash table with capacity H and a randomizing hash function has average chain length $\alpha = N/H$, where N is the current number of elements. The expected cost of insertion or lookup in separate chaining is $O(1 + \alpha)$. We keep α effectively constant by picking a suitable table size (like 40,001). Therefore, each hash operation is $O(1)$ on average.
- No other step introduces more than a constant overhead, so total is $O(1)$.

Add_jockey(jockeyId, teamId):

Algorithm:

1. If jockeyId ≤ 0 or teamId ≤ 0 , return INVALID_INPUT.
2. Look up jockeyId in the Jockey hash. If found, return FAILURE.
3. Find teamId in the Team hash. If not found, return FAILURE.
4. Allocate a new Jockey object with record = 0 and pointer to the found team.
5. Insert the jockey into the Jockey hash.
6. Return SUCCESS.

Time Complexity: All lookups/insertions are $O(1)$ average in a hash, so total is $O(1)$ expected.

Math: same separate chaining hash argument as before.

Update_match(victoriousJockeyId, losingJockeyId):

Algorithm:

1. If the IDs are invalid or the same, return INVALID_INPUT.
2. Look up each jockey in the Jockey hash. If either is missing, return FAILURE.
3. For each jockey, do Union-Find find to get the root Team. If they have the same root or if either root is inactive, return FAILURE.
4. Increase winner->record by 1, decrease loser->record by 1.
5. Remove each root from the record hash: $O(1)$.
6. rootW->record++; rootL->record--.
7. Re-insert them in the record hash with updated record.
8. Return SUCCESS.

Time Complexity:

- Two hash lookups ($O(1)$ each).
- Two union-find “find” calls. Typically union-find with path compression is near $O(1)$ amortized.
- Remove + Insert in the record hash: $O(1)$ each.

Therefore, `update_match` is $O(1)$ expected amortized.

Math:

- Union-Find with path compression has well-known upper bound $O(\alpha(n))$ per operation, where α is the inverse Ackermann function. This is < 5 for all practical n . So it is effectively constant time.
- The hash operations are again $O(1)$ expected. Summing constant or near-constant terms yields $O(1)$ total.

Merge_teams(teamId1, teamId2):

Algorithm:

1. If `teamId1` ≤ 0 or `teamId2` ≤ 0 or the IDs are the same, return `INVALID_INPUT`.
2. Look up both teams in the Team hash. If either missing, return `FAILURE`.
3. Union-Find find their roots. If the roots are the same or inactive, return `FAILURE`.
4. Compare root records. The higher record is the winner. On tie, pick `root1` as winner.
5. Remove both from the record hash.
6. Sum the records into the winner's root.
7. Mark loser as `active = false`, ban `loser->teamId`, link `loser->parent = winner`.
8. Re-insert winner with updated record in the record hash.

Time Complexity:

- 2 team hash lookups $\Rightarrow O(1)$ each.
- 2 union-find finds $\Rightarrow O(1)$ amortized.
- removal + insertion in record hash $\Rightarrow O(1)$ each.
- insertion into banned hash $\Rightarrow O(1)$.

Hence total is $O(1)$ expected.

Math:

- Again, each operation on a hash is expected $O(1)$, union-find find is $O(\alpha(n)) \sim O(1)$.
- Summation $\Rightarrow O(1)$.

Unite_by_record(record):

Algorithm:

1. If record ≤ 0 , return INVALID_INPUT.
2. In the record hash for record, find exactly one active root with that record. In the record hash for -record, find exactly one active root with -record.
3. If not exactly one in each, return FAILURE.
4. Merge them (like merge_teams): the positive-record root is the winner, negative is the loser.

Time Complexity:

- Checking each chain is $O(k)$ where k is the number of teams that currently have that record. Usually we expect a small chain or at worst $O(m)$ if many teams share the same record. But typically, the problem statement says “if exactly 2 teams with \pm record, then unite,” so it is effectively $O(1)$.
- The final merge is $O(1)$.

So unite_by_record is $O(1)$ average (amortized).

Math:

- The chain length for a hash bucket is expected α , typically a small constant if we spread out the record values well.
- Merging is $O(1)$. Summation $\Rightarrow O(1)$.

Get_jockey_record(jockeyId):

Algorithm:

1. If jockeyId ≤ 0 , return INVALID_INPUT.
2. Look up jockeyId in Jockey hash. If absent, return FAILURE.
3. Return that jockey's record field.

Time Complexity: Single hash lookup $\Rightarrow O(1)$.

Math: $E(\text{chain length}) \approx \alpha$. So $O(1)$ average.

Get_team_record(teamId):

Algorithm:

1. If teamId ≤ 0 , return INVALID_INPUT.
2. Team hash lookup. If absent, return FAILURE.
3. Union-Find “find” to get the root. If that root is not active, return FAILURE.
4. Return root->record.

Time Complexity:

- $O(1)$ hash, plus $O(1)$ union-find.
So total $O(1)$.

Math: same hashing + union-find argument.

Proving the Time Complexity requirements:

Hash-Table $O(1)$ average:

A separate-chaining hash with table size H and a uniform hash function typically yields an expected chain length $\alpha = N/H$. So an insertion or lookup is $O(1+\alpha)$ expected. If we pick H around the maximum N , α stays near 1, giving us $O(1)$ average time.

Formally, for each operation, the expected insertion or search cost is:

$E[\text{Time}] = O(1+\alpha)$, where α is constant or near-constant. Summing constants yields $O(1)$.

Union-Find $O(\alpha(n))$:

In typical union-find with path compression, the amortized cost of each union or find is $\alpha(n)$, where α is the inverse Ackermann function. For all practical n up to huge sizes, $\alpha(n) \leq 4$. This is effectively **constant** time. Even in a simplified approach, merges remain short if we always point the loser to the winner. Summation \Rightarrow near **$O(1)$** .

Space Complexity $O(n+m)$:

We must store:

1. One Team node for each team that is ever added, including those that lost merges. That is up to m .
2. One Jockey node for each jockey $\Rightarrow n$.
3. Each node is in exactly one chain for its ID or record, plus an array of size $\sim 40\,001$ for each hash. The sum of pointers and overhead remains linear in $n+m$.
4. The “banned ID” hash can have at most m entries (once per losing ID).
5. We do not keep any additional structures that grow faster than n or m .

Hence total memory is at most some constant factor times $(n+m)$.

Formally:

$$\text{MemoryUsed} = O(n) + O(m) = O(n+m).$$

Conclusion:

- The data structure is built around Union-Find for merging teams plus hash tables for quick lookups (by team ID, jockey ID, record, banned ID).
- All major functions achieve $O(1)$ expected time (amortized). The union-find “find” is near-constant due to path compression, and each hash operation is $O(1)$ on average.
- The entire system uses $O(n + m)$ memory, meeting the problem’s requirements.

Hence we satisfy the assignment’s time and space complexities with a combination of hashing and union-find.

